

Data scientist with around 5 years leading full-cycle AI and quantitative projects across investment analytics, QIS risk, and agentic AI. Works directly with the Group CEO on firm-wide AI for Nan Fung's Biotech VC and family-office arms – building autonomous learning loops and an autonomous AI investment-agents team, and shipping LLM-powered tools that influence deal screening and investment decisions. Earlier built the cross-asset risk aggregation engine at Premialab, serving leading global investment banks and institutional clients managing US\$20tn AUM.

Relevant Experience

Nan Fung International Properties – Senior Associate, Data Science

Jan 2024 – Present

- Work directly with the Group CEO on firm-wide AI for the Biotech VC and family-office arms – building autonomous learning loops and an autonomous AI investment-agents team, and driving emerging-technology direction across the group.
- Built a multi-agent platform with a centralised knowledge base over internal “skills” and documents, used across investment, leasing, and operations teams.
- Proactively built and shipped a full-cycle AI investment-screening platform that cut asset underwriting from a full day to minutes per deal – now used by the deal team, structurally changing how opportunities are sourced, screened, and pursued.
- Built a reusable agent framework with a pluggable memory system supporting all production AI agents, lifting agentic task success rates by 20% and cutting end-to-end build cycles by 70% (weeks → days).
- Owned full-cycle AI delivery: React / TypeScript front-ends over Python (FastAPI) backends with ingestion-to-reporting pipelines, deployed as human-in-the-loop AI agent SaaS on AWS EKS via GitHub Actions CI.
- Deployed NLP, forecasting, and LLM-based synthetic data generation for valuation signals, comparables, and investment-scenario stress testing.
- Applied CART analysis on campaign data to identify high-potential customer segments, driving a 4× conversion lift.

Premialab – Fintech Data Scientist

July 2022 – Jan 2024

- Led development of the cross-asset risk aggregation engine for QIS strategies, combining additive sensitivities (Greeks) and non-additive metrics (Historical VaR), serving leading global investment banks and institutional clients managing US\$20tn AUM with 60% faster reports and reliable pre-T+1 market-open risk.
- Collaborated with risk analysts and the CRO to perform risk validation, back-testing, and benchmarking for QIS derivative strategies.
- Designed and deployed a real-time anomaly detection framework to monitor data quality and trigger escalation alerts, ensuring reliability across client-facing reports.
- Rebuilt data extraction pipelines using asynchronous microservices, cutting processing time and infrastructure costs by 90%.
- Optimized SQL processes, cutting query time by 90%.

HKUST Fintech Research Project – Data Analytics Research Assistant

Aug 2021 – July 2022

- Fine-tuned GPT/BERT models on NVIDIA DGX for domain-specific text classification, lifting accuracy by 10pp.
- Led an HKMA industry report using hierarchical clustering, and a federated ML project with SHAP explainability.

Société Générale – Cross-Asset Risk Trainee

Aug 2020 – Jul 2021

- Produced daily reports analysing same-day market news and policy changes for group-wide stakeholders.
- Participated in the Data Science Lab team, independently applying Lasso Regression and PCA for global asset movements and cross-asset stress testing.
- Built Power BI dashboards cutting report time from 3 days to 10 minutes – still in use 3+ years later.
- Built real-time web apps for Twitter streaming and news scraping, leveraging BERT for sentiment analysis.

Education

Georgia Institute of Technology – MSc Analytics (Part-time), GPA 4.0 / 4.0

Jan 2024 – Present

University of Bath – MSc Computer Science, Distinction

Sept 2019 – Sept 2020

University of Bristol – BSc (Hons) Economics, Upper Second Class

Sept 2016 – Jul 2019

Technical Proficiencies

Programming – Proficient: Python, SQL; Familiar: TypeScript / JavaScript, HTML / CSS

AI / ML – LLMs (OpenAI, Anthropic, open-source), RAG, agentic workflows, LangChain, prompt & eval frameworks, PyTorch, scikit-learn, fine-tuning, NLP, forecasting, statistical modelling

Data / Cloud – AWS (EKS, S3), Azure, FastAPI, Docker, Git, GitHub Actions; ingestion-to-reporting pipelines

Languages – Fluent English, Mandarin; Native Cantonese

Achievements & Certifications

Champion (2025) – AWS AI League Hong Kong (real-time fine-tuning, AWS TechFest).

Judging Panellist (2024) – HKU Hack4SGD Hackathon.

Certifications: Azure (Fundamentals & AI), DeepLearning.AI TensorFlow Developer, IBM Data Science.